


SQL SELECT

Part I



Agenda




- Column projections and aliasing
- The WHERE clause
- The ORDER BY clause
- Table inner joins
- Table outer joins
- Multi-table and self-referencing joins

SQL SELECT: Part I

The End



SQL SELECT



SQL SELECT Statement

- The most popular and complex statement in the SQL language
- Used to query data; queries can range from the simple to the very complex
- Many different components to the statement, which are called clauses
- These clauses affect the operation of the command itself
- Most queries are intuitive

Recall: SQL SELECT



SELECT

col1, col2, ...

FROM *table*

WHERE *condition*

ORDER BY *columns*


Columns to display
(projection)

Table to use


Only return rows
matching this condition

Sort row output by
data in these columns


SELECT Statement Processing v1.0



How we write it

- 
1. TOP/DISTINCT
 2. (Projection)
 3. FROM
 4. WHERE (selection)
 5. ORDER BY

How it is processed

- 
1. FROM
 2. WHERE (selection)
 3. (Projection)
 4. ORDER BY
 5. TOP/DISTINCT

Query Processing

1. Start with this table.
67 rows × 12 columns

5. Take the first 5 rows.
5 rows × 3 columns

3. Project only these columns.
15 rows × 3 columns

```
1 select top 5
2 employee_firstname, employee_lastname, employee_pay_rate
3 from employees
4 where employee_department = 'Customer Service'
5 order by employee_pay_rate desc
```

4. Sort by this column
descending order.
15 rows × 3 columns

2. Select only rows
matching this
condition.
15 rows × 12 columns


Results			
	employee_firstname	employee_lastname	employee_pay_rate
1	Michael	Fudge	2724.5495
2	Joy	Touda-World	1813.0920
3	Amber	Wavesofgrain	1791.9112
4	Willie	Pas-D'course	1116.0667
5	Artie	Choke	20.4217

SQL SELECT

The End



SELECT Projections



Projections

- Projections determine the columns in the query output.

SELECT student_name, student_gpa **FROM** students

students			
student_id	student_name	student_gpa	student_year
1	Robin Banks	4.00	Freshman
2	Victor Edance	3.81	Junior
3	Erin Yortires	2.40	Junior
4	Phil McCup	2.75	Senior



student_name	student_gpa
Robin Banks	4.00
Victor Edance	3.81
Erin Yortires	2.40
Phil McCup	2.75

Projections (cont.)

- If we do not know the columns, we can use the asterisk (*) to project all columns in the FROM.
- Since the output is not a table, we can project the same column multiple times.
- We can project expressions, which are applied to every row. These are called derived columns.
- We can rename any column using the AS keyword, creating a column alias.

Arithmetic Operators

Operator	Purpose	For types	Example
+	Addition	Exact and approximate numeric types	$7 + 5.6 = 12.6$
-	Subtraction	Exact and approximate numeric types	$14 - 3 = 11$
*	Multiplication	Exact and approximate numeric types	$0.1 * 3 = 0.3$
/	Integer division	Integers	$14 / 4 = 3$
/	Division	Exact and approximate numeric types	$14 / 4.0 = 3.5$
%	Remainder	Integers	$14 \% 4 = 2$
+	Concatenation	String types	'Mi' + 'ke' = 'Mike'

Data Type Casting



Change the data type of an expression



Useful for type conversion, truncation,
parsing



CAST(*expression AS new_type*)


SELECT Projections

The End



Demo

SELECT Projections




Demo: SELECT Projections




- We will use the payroll database.
- Selecting specific columns
- Column concatenation
- Column aliasing with AS
- CAST() data to different types
- The need for CAST(): conversion errors

Demo: SELECT Projections

The End



SELECT WHERE



Selections (WHERE Clause)

- Selections determine the rows in the query output.

```
SELECT * FROM students WHERE student_year = 'Junior'
```

students			
student_id	student_name	student_gpa	student_year
1	Robin Banks	4.00	Freshman
2	Victor Edance	3.81	Junior
3	Erin Yortires	2.40	Junior
4	Phil McCup	2.75	Senior



student_id	student_name	student_gpa	student_year
2	Victor Edance	3.81	Junior
3	Erin Yortires	2.40	Junior

Selections (WHERE Clause) (cont.)

- If the WHERE clause is omitted, all rows are returned.
- The condition included in the WHERE clause is a Boolean expression.
- If the Boolean expression is true, the row is included in the output.
- Relational operators are Boolean expressions that compare values.
- Logical operators are Boolean expressions to connect relational expressions.

Relational Operators

Operator	Definition	Example of use
=	Equal to	<code>1 = 1</code>
!=	Not equal to	<code>'Mike' != 'Fudge'</code>
>	Greater than	<code>3 > 1</code>
<	Less than	<code>1 < 3</code>
>=	Greater than or equal to	<code>4 >= 4</code>
<=	Less than or equal to	<code>6 <= 10</code>
In	Set inclusion	<code>4 in (2,4,6,8)</code>
Not in	Set exclusion	<code>'eggs' not in ('bacon', 'cheese', 'toast')</code>
Is null	Column null check	<code>Column_name is null</code>
Is not null	Column not null check	<code>Column_name is not null</code>
Between	Inclusive numeric range	<code>Column_name between 0 and 4</code>
Like	Text pattern matching	<code>'Michael' like 'Mi%'</code>
Not Like	Text pattern not matching	<code>'Michael' not like 'Mike%'</code>

Logical Operators

Operator	Definition	Example of use
AND	Evaluates to true only when both expressions are true, otherwise false	Customers with a gmail.com email address in the 315 area code <code>customer_phone like '315%' AND customer_email like '%@gmail.com'</code>
OR	Evaluates to false only when both expressions are false, otherwise true	Students on academic warning below 2.0 GPA or on the Dean's list above 3.4 GPA <code>student_GPA >= 3.4 OR student_gpa < 2.0</code>
NOT	Negation; false becomes true, true becomes false	Employees not in the 'Toys' department <code>NOT Employee_department = 'Toys'</code> Note—this is the equivalent to: <code>Employee_department != 'Toys'</code>


SELECT WHERE

The End



Demo

SELECT With WHERE



Demo: SELECT With WHERE




- We will use the payroll database
- WHERE clauses
- Relational operators \geq , $<$
- Pattern matching with Like % and Like?
- Logical operators, AND/OR

Demo: SELECT With WHERE

The End



SELECT CASE





SELECT

CASE

CASE Statement

- Tests a sequence of Boolean expressions
- The first one true returns the value
- Can be added anywhere a value is yielded—for example, projection, selection

CASE

WHEN *condition1* THEN *value1*

WHEN *condition2* THEN *value2*

...

[ELSE *catch_all_value*]

END


SELECT CASE

The End



Demo

CASE




Demo: SELECT With CASE




- We will use the payroll database
- CASE to create pay bands for sales associates
- Why doesn't this work in the WHERE clause?
- Workarounds for this?

Demo: CASE

The End



SELECT ORDER BY



Sorting Output (ORDER BY Clause)

- The ORDER BY clause is used to sort the query output.
- You must specify the columns by which you will sort.
- The ASC keyword sorts in ascending order and is the default.
- The DESC keyword sorts in descending order.
- If the ORDER BY clause is omitted, rows are returned in the default table order based on the primary key.

TOP and DISTINCT Keywords

TOP n

- Limits output to the first n rows
- This is based on the sorted output of the query

DISTINCT

- Removes duplicate rows from the query output
- Guarantees the output has entity integrity


SELECT ORDER BY

The End



Demo

TOP, DISTINCT, ORDER BY



Demo: TOP, DISTINCT, ORDER BY




- We will use the payroll database
- DISTINCT Example
- ORDER BY to arrange output
- Combined with TOP to produce a specific number of rows, such as the highest or lowest

Demo: TOP, DISTINCT, ORDER BY

The End




SELECT Joins



Joins: Inner Join

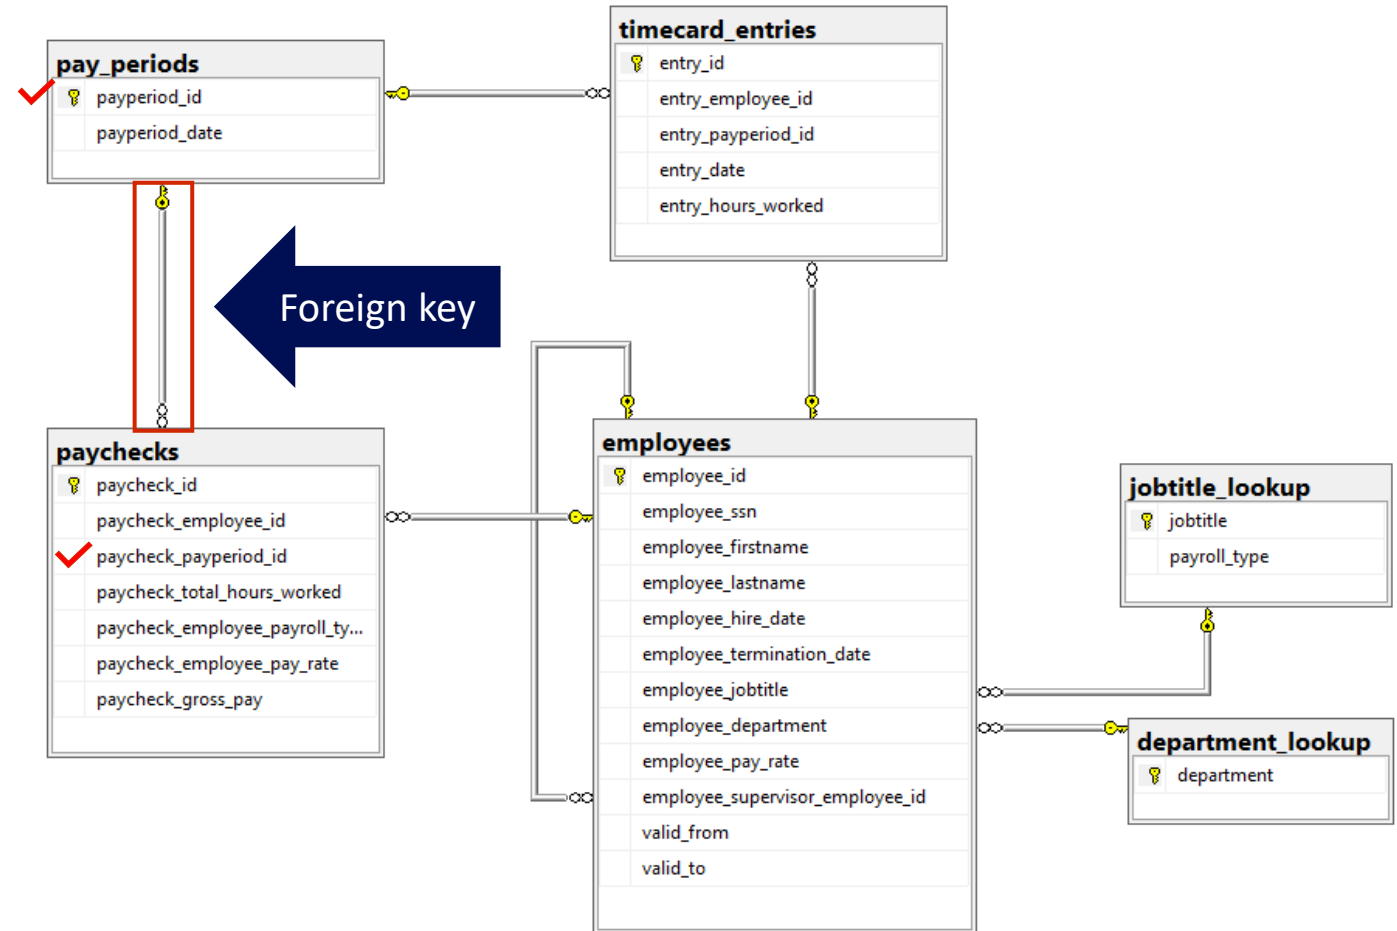
```
FROM table_a  
    JOIN table_b ON column_a = column_b
```

- Processed first in the SELECT statement
- Columns must be the same data type
- Typically PK–FK pairs
- Output is a temporary table in the query pipeline consisting of matching rows



Payroll Database Internal Model

```
FROM paychecks
JOIN pay_periods
ON payperiod_id =
paycheck_payperiod_id
```




SELECT Joins
The End



Demo

Inner Join



Demo: Inner Join



- We will use the demo and payroll databases.
- Inspect bbplayers and bbteams tables.
- Implement inner join.
- Join output is a subset of both tables.
- Note that you can join any column of the same data type, but that does not mean it is sensible to do so.

Demo: Inner Join

The End




Outer Joins



Joins: Left Outer Join

```
FROM table_a LEFT JOIN table_b  
    ON column_a = column_b
```


- All rows in *table_a* are included, including those matching the ON condition.



Joins: Right Outer Join

FROM *table_a* RIGHT JOIN *table_b*
ON *column_a* = *column_b*


- All rows in *table_b* are included, including those matching the ON condition.
- LEFT and RIGHT determine which table participates in the outer join.



Joins: Full Outer Join

```
FROM table_a FULL JOIN table_b  
    ON column_a = column_b
```

- All rows in *table_b* are included, including those matching the ON condition
- Like performing a LEFT and RIGHT join simultaneously




Outer Joins
The End



Demo

Outer Joins



Demo: Outer Joins



- We will use the demo and payroll databases
- Players not on a team
- Teams without players
- Both!
- Null checks in the query to get exceptions, such as only players without teams
- Which employees have not filled out a timecard?

Demo: Outer Joins

The End



Advanced Joins



More Joins

- It is not limited to two table joins. We can join multiple tables.

```
SELECT *  
FROM table_a  
    JOIN table_b ON table_b_col = table_a_col  
    JOIN table_c ON table_c_col = table_a_or_b_col
```

- We can alias tables to simplify complex joins.

```
SELECT a.*  
FROM table_a AS a  
    JOIN table_b ON table_b_col = a.table_a_col
```

- We can use table aliasing to join a table to itself!


Advanced Joins

The End



Demo

Multi-Table Joins and Self-Joins



Demo: Multi-Table Joins and Self-Joins




- We will use the payroll database
- Joining three tables in a complex join
- Table aliasing in a complex join
- Joining a table to itself; employees and supervisors

Demo: Multi-Table Joins and Self-Joins

The End



Summary



Summary



- The processing order of the SELECT statement differs from how it is written.
- Projections filter columns.
- Selections filter rows with the WHERE clause.
- The CASE statement is used for conditional evaluation.
- Joins combine tables based on matching columns.
- Outer joins allow for including rows in a table that do not match the inner join output.

Summary

The End

