

Transaction Management and Concurrency Control



Agenda




- What are database transactions?
- What are the ACID properties of a transaction?
- How to write transaction-safe SQL code.
- What is concurrency control, and how is it related to transactions?
- Understand advanced concurrency issues like versioning, locking, and deadlocks.

Transaction Management and
Concurrency Control

The End




What Are Transactions?



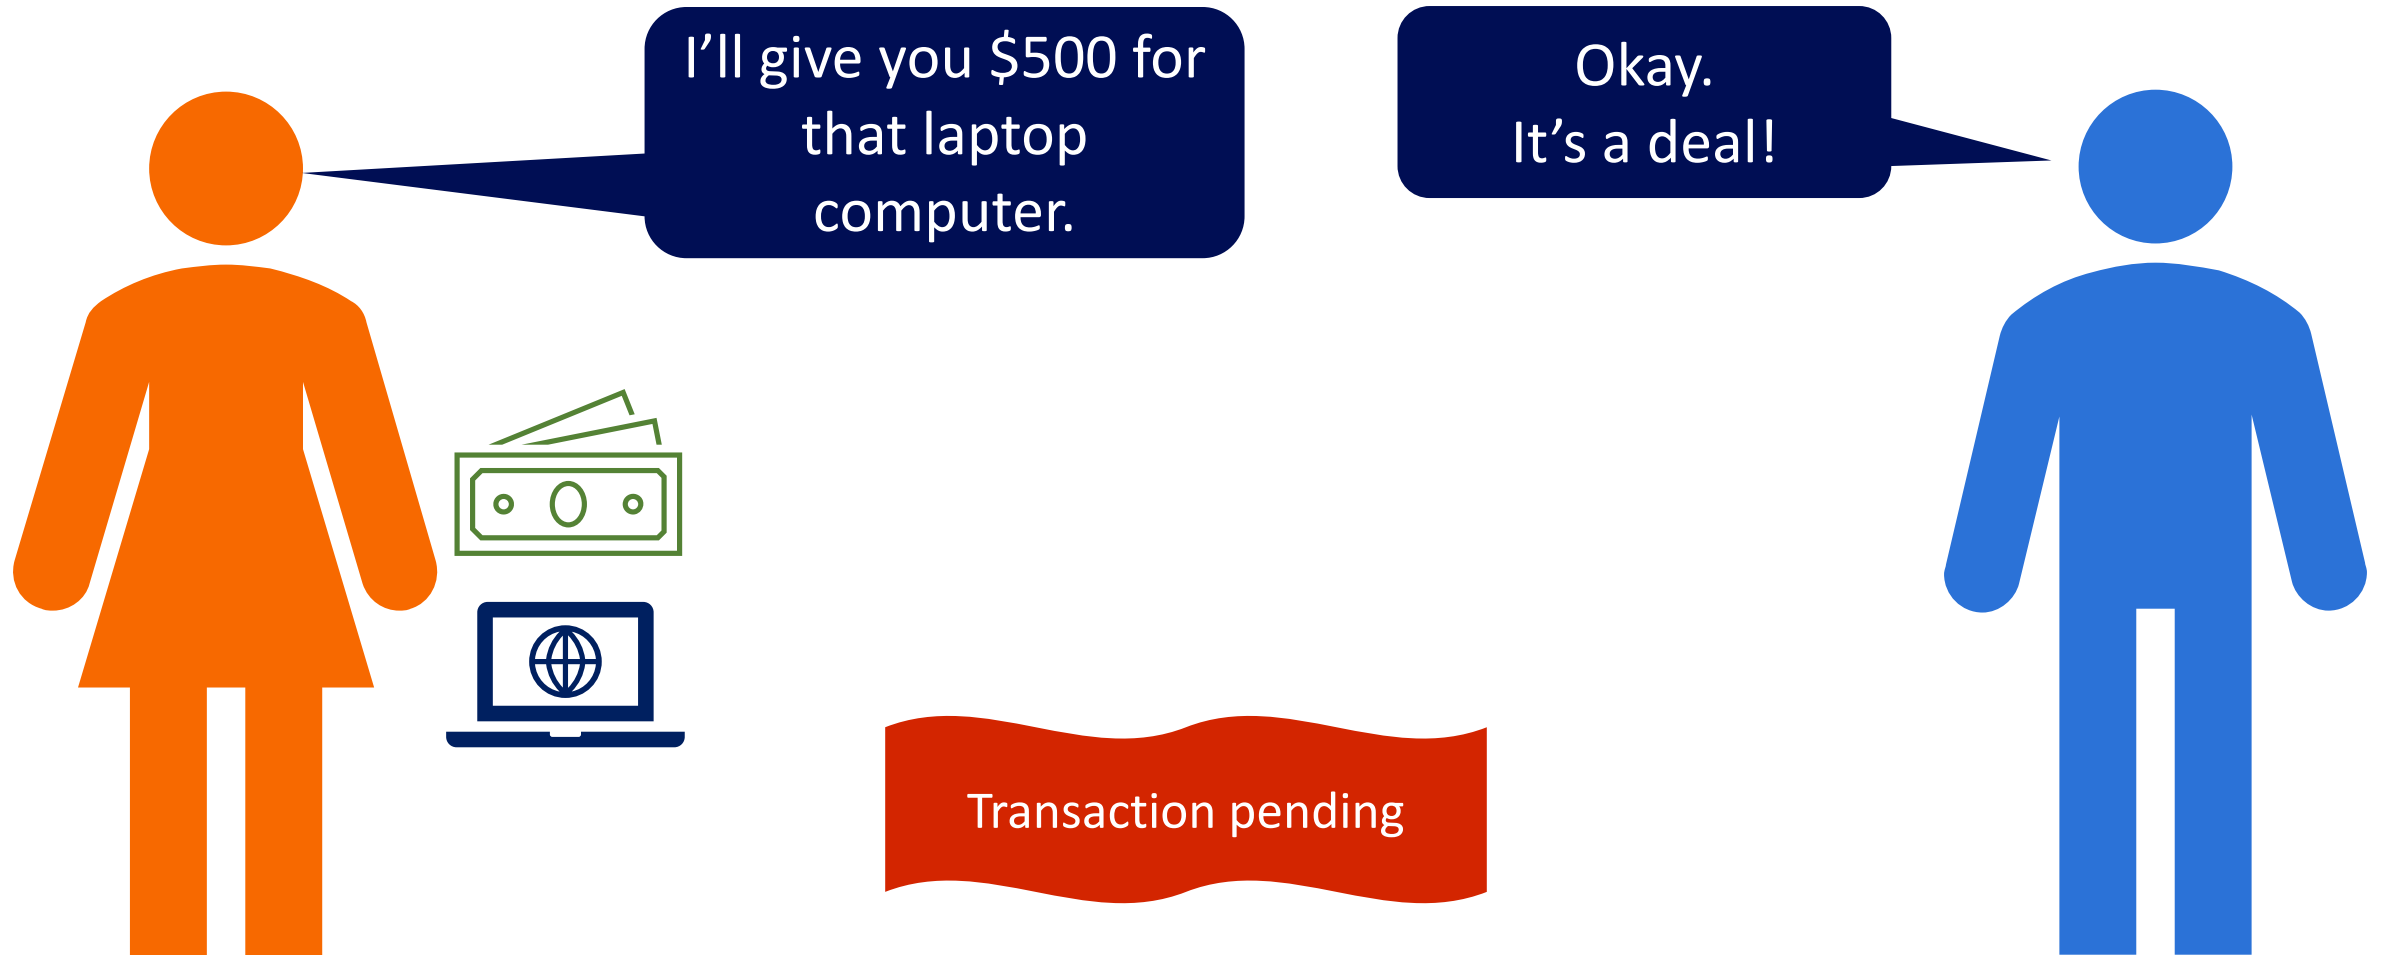
Transaction

- Any logical unit of work in your database management system, read, write, or combination thereof
- It is data logic
- Typically consists of several read/write operations
- Must succeed or fail as a whole
- Operates independently from other units of work
- Doesn't know about anything other than itself

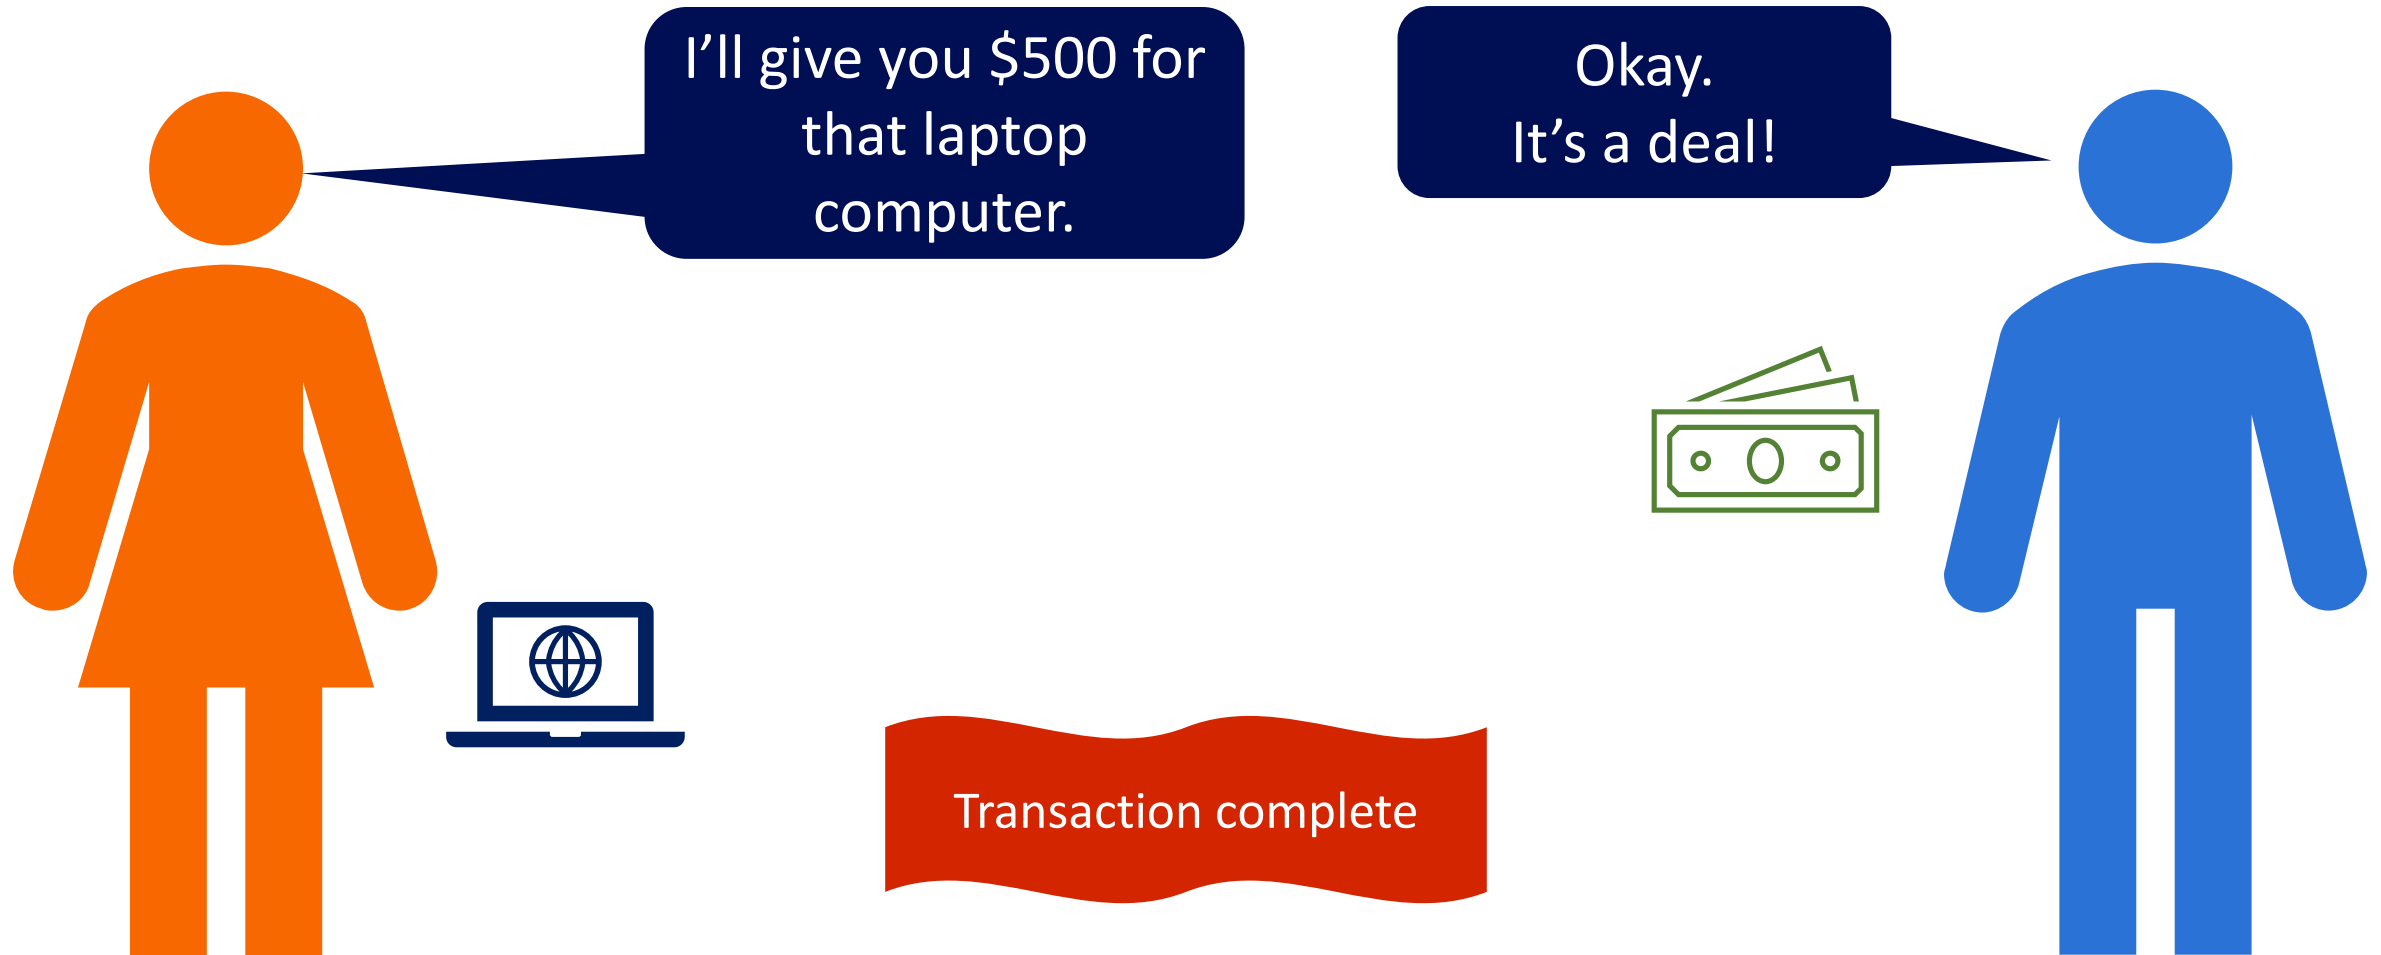
Real-Life Example



Real-Life Example




Real-Life Example




Database Transactions

Guarantee No Intermediary States




What Are Transactions?

The End



Demo

The Need for Transactions



Demo: The Need for Transactions




- We will use the Azure Data Studio application
- We will use the demo database
- Accounts table
- Understand the need for transactions and the problem of intermediary states
- This code does not use transactions, but what's the problem?
 - Check constraint firing
 - No rows affected

Demo: The Need for Transactions

The End



ACID Properties



Transaction ACID Properties

- Atomic
 - Transaction cannot be subdivided; logical unit of work; completes as a single unit or not at all
- Consistent
 - Transitions data from one state to another, constraints intact
- Isolated
 - Database changes not revealed to other users until after transaction has completed
- Durable
 - Database changes are permanent once committed

Example: Transfer Funds

- Transfer \$1,000 from savings to checking

Checking \$400

Savings \$1,500

Before

Example: Transfer Funds (cont.)

- Transfer \$1,000 from savings to checking

Checking \$1,400

Savings \$500

After

Atomic: Transfer Funds

- Transfer \$1,000 from savings to checking
 - Step 1 remove \$1,000 from savings
 - Step 2 add \$1,000 to checking

Checking \$400

Savings \$500

Nobody ever sees this step in the process since the transaction is treated as a single operation.

Consistent: Transfer Funds

- Transfer \$2,000 from savings to checking
- Business rule accounts: 0 or more

Checking \$400

Savings \$-500

This transaction cannot be completed. Doing so would leave the data in an inconsistent state.

Consistent: Transfer Funds (cont.)

- Transfer \$2,000 from savings to checking
- Business rule accounts: 0 or more

Checking \$400

Savings \$1,500

Because of this inconsistency, the data are reverted to their original state.

Isolated: Transfer Funds

- Transfer \$1,000 from savings to checking

Checking \$400

Savings \$1,500

Someone else on the account tries to read the savings balance at this point in time.
They wait...

Isolated: Transfer Funds (cont.)

- Transfer \$1,000 from savings to checking

Checking \$1,400

Savings \$500

... until the transaction is completed, then they see the savings balance of \$500.


Durable: Transfer Funds

- Transfer \$1,000 from savings to checking

Checking \$400

Savings \$1,500

The state of the data
before...



Durable: Transfer Funds (cont.)


- Transfer \$1,000 from savings to checking

Checking \$1,400

Savings \$500

The state of the data
before...

... and after are
stored in the
database
permanently.




ACID Properties

The End



Transaction Support in SQL Server



Transactions in SQL Server

- We simply take the code we have:

```
BEGIN TRANSACTION
```

```
UPDATE accounts SET checking ...
```

```
UPDATE accounts SET checking ...
```

```
COMMIT TRANSACTION
```

- And tell SQL Server to group them as a transaction—sort of

T-SQL Transaction Commands

- **BEGIN TRAN[SACTION]**
 - Starts the marking point of a transaction; any data manipulation statements after this point are not durable
- **COMMIT [TRAN[SACTION]]**
 - Marks the end of a successful transaction; at this point, all data changes to become durable
- **ROLLBACK [TRAN[SACTION]]**
 - Undoes any data manipulation statements since the beginning of the transaction
- **@@TRANCOUNT**
 - A T-SQL Internal variable that identifies pending transactions


Transaction Support in SQL Server

The End



Demo

SQL Transactions



Demo: SQL Transactions



- We will use the Azure Data Studio application
- We will use the demo database
- Without BEGIN TRANSACTION, all commands are durable
- With BEGIN TRANSACTION, they are not!
- @@TRANCOUNT
- COMMIT and ROLLBACK
- Transactions are isolated, consistent, atomic, and durable

Demo: SQL Transactions

The End



Transaction Safe Code



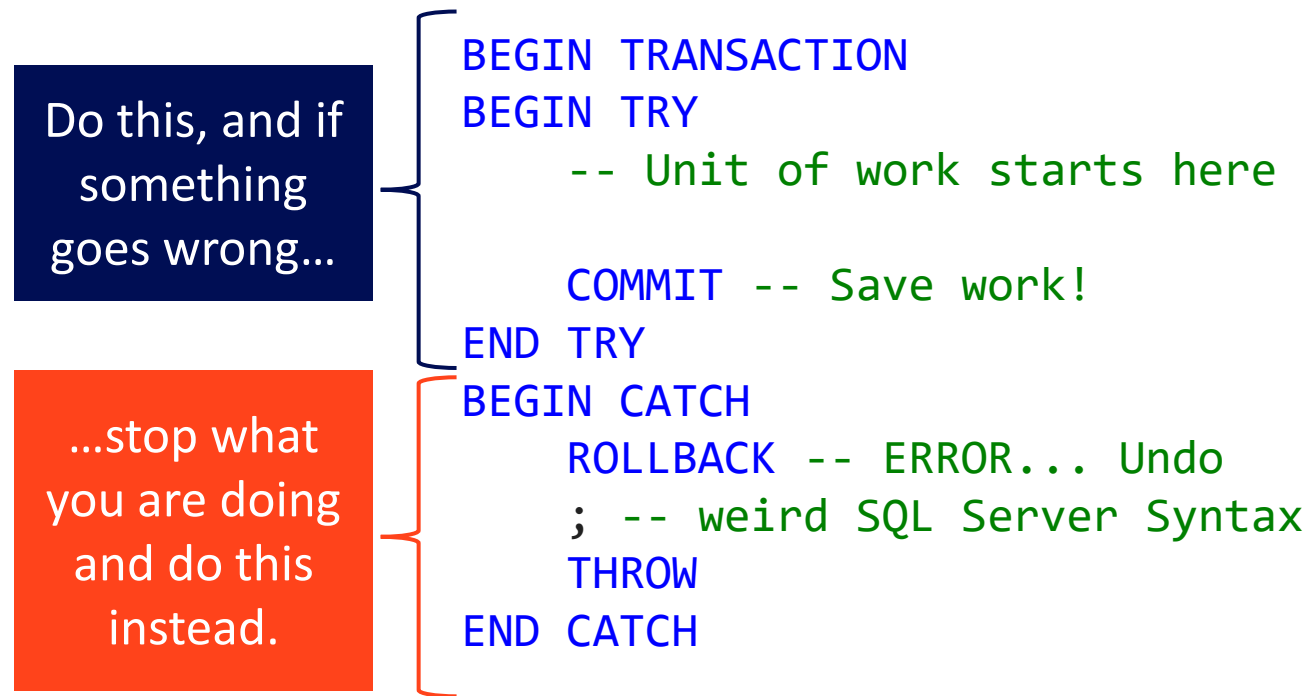
Transaction Safe Code 101: How Do You Know When to Rollback?

We want to rollback when:

- An error occurs
 - Database cannot write to the table (disk is full)
 - There is a physical hardware/operating system issue
 - A database constraint is violated (check, PK, FK, etc.), which will result in an inconsistency
- There is custom data logic
 - The expected number of rows are not affected, for example

SQL Error Handling: TRY... CATCH

The SQL TRY... CATCH statement makes transaction management of errors simple.




Transaction Safe Code

The End



Demo

Transactions TRY/CATCH



Demo: Transactions TRY/CATCH




- We will use the Azure Data Studio application.
- We will use the demo database.
- Re-write our p_transfer_funds to use transactions.
- Test it out:
 - Check constraint firing
 - No rows affected
- Not perfect, but better!

Demo: Transactions TRY/CATCH

The End



Rollback on Custom Data Logic



Rollback From Custom Data Logic

- We might want to rollback based on custom data logic
- For example, when you update a row but no rows are affected, how do you handle this?

@@ROWCOUNT—reports number of rows affected from an SQL statement

```
SELECT * FROM accounts  
(2 rows affected)  
PRINT @@ROWCOUNT  
2
```

SQL THROW a Custom Error

Use the SQL THROW statement to launch a custom user-defined error

THROW error_number, message, state

- error_number — a customer number ≥ 50000
- Message — a string error message
- State — a number 0–255 for message state, usually a 1


Rollback on Custom Data Logic

The End



Demo

Rollback On Custom Data Logic



Demo: Rollback on Custom Data Logic




- We will use the Azure Data Studio application.
- We will use the demo database.
- Re-write our p_transfer_funds to use transactions.
- Handle custom business logic for rows affected.

Demo: Rollback Custom Data Logic

The End




Concurrency Control




Concurrency Control

- Problem: in a multiuser environment, simultaneous access to data can result in interference and data loss
- Solution: concurrency control
 - The process of managing simultaneous transactions against a database so that data integrity is maintained, and the operations do not interfere with each other in a multi-user environment.
 - Concurrency control helps keep transactions isolated.
 - Concurrency control is a part of every modern DBMS.

Example: Concurrency Question



```
select * from students  
| where student_year_name = 'Freshman'
```



```
update students set student_firstname = 'Robyn'  
| where student_id = 1  
  
update students set student_firstname = 'Bucky'  
| where student_id = 16
```



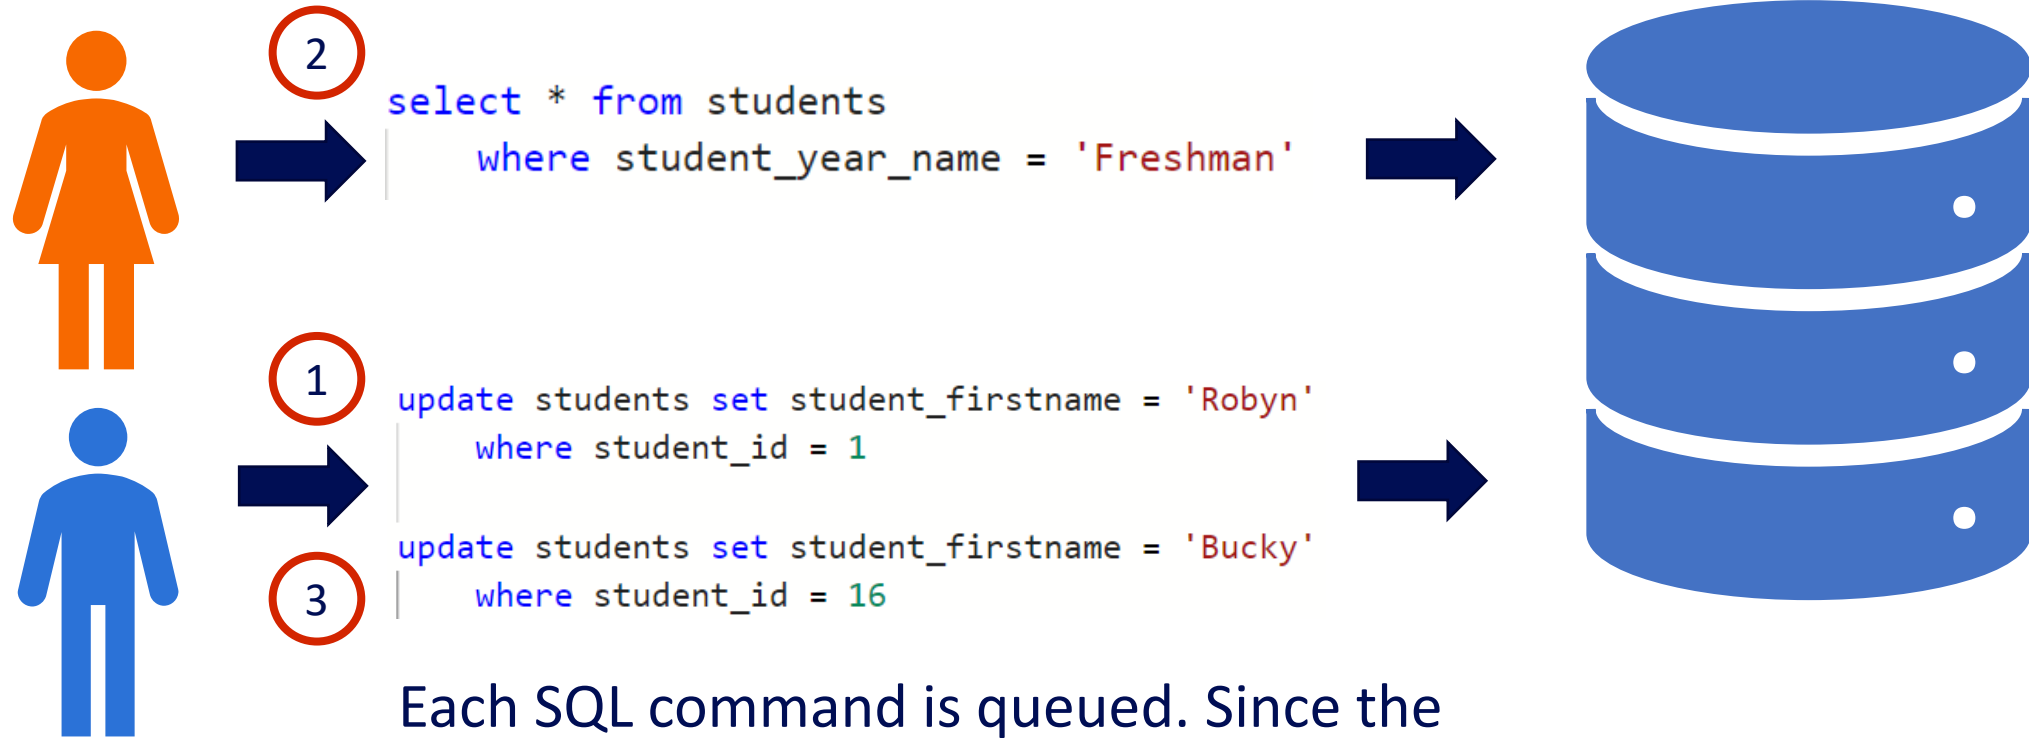
What does the top person see as output?

Robin Robyn Robyn

Buck Buck Bucky

student_id	student_firstname	student_lastname	student_year_name	student_gpa
1	Robin	Banks	Freshman	4.000
2	Victor	Edance	Freshman	2.404
8	Lola	Dabridgeda	Freshman	2.732
10	Phil	McCup	Freshman	2.705
16	Buck	Naked	Freshman	2.434
21	Cook	Myefoud	Freshman	3.593
25	Oliver	Stuffission	Freshman	3.118

Example: Concurrency Answer



Each SQL command is queued. Since the bottom person did not use a transaction, they could get ordered:


Robyn
Bucky

Concurrency Control

The End



DBMS Without Concurrency Control



Lost Update Problem

Occurs when one update overwrites another

Transaction A	Time	Transaction B
READ Balance = \$500	1	---
---	2	READ Balance = \$500
---	3	Withdrawal \$300 WRITE Balance = \$200
Withdrawal \$100 WRITE Balance = \$400	4	---
	5	READ Balance = \$400?!

Balance should be \$100, but it's \$400!

Dirty/Inconsistent Read Problem

Occurs when a transaction reads uncommitted data

Transaction A	Time	Transaction B
---	1	READ City = 'Utica'
---	2	UPDATE City = 'Rome'
READ City = 'Rome'	3	---
---	4	FAIL: Rollback


A's city = 'Rome' and B's city = 'Utica'

DBMS Without Concurrency Control

The End



Serializability and Locking



How Do DBMS Maintain Concurrency Control?

- Serializability
 - Finish one transaction before starting another
- Locking mechanisms
 - The most common way of achieving serialization
 - Data that are retrieved for the purpose of updating are locked for the updater
 - No other user can perform a write operation until unlocked

Serializability and Locking

This demonstrates an exclusive lock.

Transaction A	Time	Transaction B
Request account, lock acquired	1	---
READ Balance = \$500	2	Request account, waiting for lock release
---	3	---
Withdrawal \$100 WRITE Balance = \$400	4	---
Lock released	5	---
---	6	Lock acquired
---	7	READ Balance = \$400

Locking Mechanisms


- Locking granularity
 - Database: used during database updates, ALTER DATABASE
 - Table: used for bulk updates, or ALTER TABLE
 - Block or page: very commonly used
 - Record: only requested row; fairly commonly used
 - Field: requires significant overhead; impractical
- Types of locks
 - Shared lock: read but no update permitted; used when just reading to prevent another user from placing an exclusive lock on the record
 - Exclusive lock: no access permitted; used when preparing to update

Serializability and Locking

The End



Versioning and Isolation Levels



Versioning

- DBMS maintains multiple versions of the data to be modified as part of the transaction.
- Concurrent reads are allowed to the data.
- Attempts to write are allowed, but subsequent writes are rolled back and restarted.

Versioning Example

Transaction A	Time	Transaction B
READ Balance = \$500	1	---
---	2	READ Balance = \$500
Withdrawal \$100 Begin transaction WRITE Balance = \$400	3	---
---	4	Withdrawal \$300 Begin transaction WRITE Balance = \$200
Commit	5	Other version pending, rollback
---	6	Restart transaction

Isolation Levels in SQL Server

- We change concurrency control by setting an isolation level in the database
- Read uncommitted: any transaction can read any uncommitted data; no locks, no concurrency control.
- Read committed: no transaction can read uncommitted data; default
- Serializable: like read committed but also locks related data across foreign keys
- Read committed snapshot: uses versioning so that there are no locks on the read operation


Versioning and Isolation Levels

The End



Demo

Concurrency Control



Demo: Concurrency Control




- We will use the Azure Data Studio application.
- We will use the demo database.
- Let's demonstrate a dirty read by playing with the isolation levels.
 - Update a row
 - Witness a lock on the table read committed
 - No lock with read uncommitted—but you read pending transaction data
 - When transaction is rolled back—yikes!

Demo: Concurrency Control

The End



Deadlocks



Deadlocks

- A deadlock situation occurs when two or more transactions are waiting for each other to give up locks
- Row level locks help mitigate this issue, but there is still a possibility
- Set the lock timeout in milliseconds
- **SET LOCK_TIMEOUT 5000—five seconds**
- Read it with **@@LOCK_TIMEOUT**
- Default is -1 (never)

Deadlock Example

Transaction A	Time	Transaction B
Begin transaction	1	---
---	2	Begin transaction
Update users where id = 1	3	---
---	4	Update blogs where id = 7
Update blogs where id = 7 (Locked by B, waiting...)	5	---
---	6	Update users where id = 1 (Locked by A, waiting...)
---	7	---


Deadlocks

The End



Demo

Deadlock



Demo: Deadlock




- We will use the Azure Data Studio application.
- We will use the tinyu database.
- Let's demonstrate a deadlock scenario.
 - Transaction 1: Update row in table A, then row in table B.
 - Transaction 2: Update row in table B, then row in table A.

Demo: Deadlock

The End



Summary



Summary



- Transactions are multiple SQL statements treated as a single unit of work.
- Transactions are atomic, consistent, isolated, and durable.
- To make SQL transaction safe, you must rollback on error, or when the expected data logic results do not match the actual results.
- Concurrency control is how multiple transactions are managed within the DBMS.
- Deadlocks occur when two or more transactions are waiting to give up locks.

Summary

The End

